

Neural Networks and deep learning – ICP4

KARTIK ANUMOLU

Student ID : 700758573

Github link:<https://github.com/kartikanumolu1/neural-assignment4.git>

Video link:

https://drive.google.com/file/d/1J4_9327I_rSTf4idXSEUptxCaUVuWUdM/view?usp=drive_link

1. Follow the instruction below and then report how the performance changed.(apply all at once)
 - Convolutional input layer, 32 feature maps with a size of 3×3 and a rectifier activation function.
 - Dropout layer at 20%.
 - Convolutional layer, 32 feature maps with a size of 3×3 and a rectifier activation function.
 - Max Pool layer with size 2×2 .
 - Convolutional layer, 64 feature maps with a size of 3×3 and a rectifier activation function.
 - Dropout layer at 20%.
 - Convolutional layer, 64 feature maps with a size of 3×3 and a rectifier activation function.
 - Max Pool layer with size 2×2 .
 - Convolutional layer, 128 feature maps with a size of 3×3 and a rectifier activation function.
 - Dropout layer at 20%.
 - Convolutional layer, 128 feature maps with a size of 3×3 and a rectifier activation function.
 - Max Pool layer with size 2×2 .
 - Flatten layer.
 - Dropout layer at 20%.
 - Fully connected layer with 1024 units and a rectifier activation function.
 - Dropout layer at 20%.
 - Fully connected layer with 512 units and a rectifier activation function.
 - Dropout layer at 20%.
 - Fully connected output layer with 10 units and a Softmax activation function Did the performance change?
2. Predict the first 4 images of the test data using the above model. Then, compare with the actual label for those 4 images to check whether or not the model has predicted correctly.
3. Visualize Loss and Accuracy using the history object

Output:

The image displays two sequential screenshots of a Google Colab notebook titled 'icp4_neural.ipynb'. The first screenshot shows the initial setup code, including imports for TensorFlow, Keras, and NumPy, setting a random seed, loading CIFAR-10 data, and normalizing it. The second screenshot shows the creation of a deep neural network model with multiple layers of convolution, pooling, and dense connections.

```
import numpy as np
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, Flatten, Conv2D, MaxPooling2D
from tensorflow.keras.constraints import MaxNorm
from tensorflow.keras.optimizers import SGD
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.optimizers.schedules import ExponentialDecay
import matplotlib.pyplot as plt

# Fix random seed for reproducibility
seed = 7
np.random.seed(seed)

# Load data
(X_train, y_train), (X_test, y_test) = cifar10.load_data()

# Normalize inputs from 0-255 to 0.0-1.0
X_train = X_train.astype('float32') / 255.0
X_test = X_test.astype('float32') / 255.0

# One hot encode outputs
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
num_classes = y_test.shape[1]

# Create the model
model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape=(32, 32, 3), padding='same', activation='relu', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(32, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Flatten())
model.add(Dropout(0.2))
model.add(Dense(1024, activation='relu', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu', kernel_constraint=MaxNorm(3)))
```

Content | Bb 15426092 | icp4_neural.ipynb | Home - Google D | icp4_neural_assign | kartikanumolu1/ne | +

colab.research.google.com/drive/10i9VaS4rb2y9eOIT7K0qODarhXMLZBjm#scrollTo=iC4fly1T-W9t

icp4_neural.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Code + Text

```
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))

# Compile model
epochs = 5
lr_rate = 0.01
lr_schedule = ExponentialDecay(
    initial_learning_rate=lr_rate,
    decay_steps=epochs * len(X_train) // 32,
    decay_rate=0.1
)
sgd = SGD(learning_rate=lr_schedule, momentum=0.9, nesterov=False)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
print(model.summary())

# Fit the model
history = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=32)

# Final evaluation of the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1] * 100))

# Predict the first 4 images of the test data
predictions = model.predict(X_test[:4])
```

Executing (6m 52s)

Content | Bb 15426092 | icp4_neural.ipynb | Home - Google D | icp4_neural_assign | kartikanumolu1/ne | +

colab.research.google.com/drive/10i9VaS4rb2y9eOIT7K0qODarhXMLZBjm#scrollTo=iC4fly1T-W9t

icp4_neural.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Code + Text

```
predicted_classes = np.argmax(predictions, axis=1)
actual_classes = np.argmax(y_test[:4], axis=1)

# Print the predictions and actual labels
print("Predicted classes: ", predicted_classes)
print("Actual classes: ", actual_classes)

# Plot the first 4 test images, predicted labels, and actual labels
fig, axes = plt.subplots(1, 4, figsize=(15, 3))
for i in range(4):
    axes[i].imshow(X_test[i])
    axes[i].set_title(f"Pred: {predicted_classes[i]}, Actual: {actual_classes[i]}")
    axes[i].axis('off')
plt.show()

# Visualize Loss and Accuracy
plt.figure(figsize=(12, 4))

# Plot Loss
plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label='train_loss')
plt.plot(history.history['val_loss'], label='val_loss')
plt.title('Model Loss')
plt.ylabel('Loss')
```

Executing (7m 1s)

Content x Bb 15426092 x icp4_neural.ipynb x Home - Google D... x icp4_neural_assign... x kartikanumolu1/ne... x + -

colab.research.google.com/drive/10i9VaS4rb2y9eOIT7K0qODarhXMLZBjm#scrollTo=iC4fly1T-W9t

icp4_neural.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Comment Share Gemini

+ Code + Text

```
plt.xlabel('Epoch')
plt.legend(loc='upper right')

# Plot Accuracy
plt.subplot(1, 2, 2)
plt.plot(history.history['accuracy'], label='train_accuracy')
plt.plot(history.history['val_accuracy'], label='val_accuracy')
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(loc='upper left')

plt.show()
for i in range(4):
    axes[i].imshow(X_test[i])
    axes[i].set_title(f"Pred: {predicted_classes[i]}, Actual: {actual_classes[i]}")
    axes[i].axis('off')
plt.show()

# Visualize Loss and Accuracy
plt.figure(figsize=(12, 4))

# Plot Loss
plt.subplot(1, 2, 1)
```

Executing (7m 6s)

24°C Partly cloudy

Search

ENG IN 10:56 11-07-2024

Content x Bb 15426092 x icp4_neural.ipynb x Home - Google D... x icp4_neural_assign... x kartikanumolu1/ne... x + -

colab.research.google.com/drive/10i9VaS4rb2y9eOIT7K0qODarhXMLZBjm#scrollTo=iC4fly1T-W9t

icp4_neural.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Comment Share Gemini

+ Code + Text

```
plt.figure(figsize=(12, 4))

# Plot Loss
plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label='train_loss')
plt.plot(history.history['val_loss'], label='val_loss')
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(loc='upper right')

# Plot Accuracy
plt.subplot(1, 2, 2)
plt.plot(history.history['accuracy'], label='train_accuracy')
plt.plot(history.history['val_accuracy'], label='val_accuracy')
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(loc='upper left')

plt.show()

Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
--------------	--------------	---------

Executing (7m 24s)

24°C Partly cloudy

Search

ENG IN 10:57 11-07-2024

colab.research.google.com/drive/10i9VaS4rb2y9eOIT7K0qODarhXMLZBjm#scrollTo=iC4fly1T-W9t

icp4_neural.ipynb

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_6 (Conv2D)	(None, 32, 32, 32)	896
dropout_6 (Dropout)	(None, 32, 32, 32)	0
conv2d_7 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d_3 (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_8 (Conv2D)	(None, 16, 16, 64)	18496
dropout_7 (Dropout)	(None, 16, 16, 64)	0
conv2d_9 (Conv2D)	(None, 16, 16, 64)	36928
max_pooling2d_4 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_10 (Conv2D)	(None, 8, 8, 128)	73856
dropout_8 (Dropout)	(None, 8, 8, 128)	0

Executing (7m 30s)

24°C Partly cloudy

colab.research.google.com/drive/10i9VaS4rb2y9eOIT7K0qODarhXMLZBjm#scrollTo=iC4fly1T-W9t

icp4_neural.ipynb

conv2d_11 (Conv2D)	(None, 8, 8, 128)	147584
max_pooling2d_5 (MaxPooling2D)	(None, 4, 4, 128)	0
flatten_1 (Flatten)	(None, 2048)	0
dropout_9 (Dropout)	(None, 2048)	0
dense_3 (Dense)	(None, 1024)	2098176
dropout_10 (Dropout)	(None, 1024)	0
dense_4 (Dense)	(None, 512)	524800
dropout_11 (Dropout)	(None, 512)	0
dense_5 (Dense)	(None, 10)	5130

Total params: 2915114 (11.12 MB)
Trainable params: 2915114 (11.12 MB)
Non-trainable params: 0 (0.00 Byte)

None

Executing (7m 39s)

24°C Partly cloudy

Content x 15426092 x icp4_neural.ipynb x Home - Google Dr x icp4_neural_assign x kartikanumolu1/h x +

colab.research.google.com/drive/10i9VaS4rb2y9eOIT7K0qODarhXMLZBjm#scrollTo=iC4fly1T-W9t

icp4_neural.ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

✓ 1m

Epoch 1/5
1563/1563 [=====] - 16s 9ms/step - loss: 1.8508 - accuracy: 0.3115 - val_loss: 1.5590 - val_accuracy: 0.5526
Epoch 2/5
1563/1563 [=====] - 16s 10ms/step - loss: 1.4047 - accuracy: 0.4875 - val_loss: 1.2387 - val_accuracy: 0.5526
Epoch 3/5
1563/1563 [=====] - 13s 8ms/step - loss: 1.1951 - accuracy: 0.5733 - val_loss: 1.1510 - val_accuracy: 0.5880
Epoch 4/5
1563/1563 [=====] - 13s 8ms/step - loss: 1.0507 - accuracy: 0.6233 - val_loss: 0.9887 - val_accuracy: 0.6514
Epoch 5/5
1563/1563 [=====] - 15s 9ms/step - loss: 0.9474 - accuracy: 0.6632 - val_loss: 0.9493 - val_accuracy: 0.6669
Accuracy: 66.69%
1/1 [=====] - 0s 141ms/step
Predicted classes: [3 8 8 0]
Actual classes: [3 8 8 0]

Pred: 3, Actual: 3

Pred: 8, Actual: 8

Pred: 8, Actual: 8

Pred: 0, Actual: 0

Executing (7m 45s)

SENSEX -0.03%

Search

ENG IN

10:57 11-07-2024

